

**ECE 763 Computer Vision**

**Project-2 Report**

**Priya Diwakar**

**Student ID 200205361**

**Face image classification**

**using a simple neural network to get familiar with steps of  
training neural networks**

## TABLE OF CONTENTS

<b>DATASET.....</b>	<b>3</b>
TRAINING DATASET .....	3
TESTING DATASET .....	3
VALIDATION DATASET .....	3
<b>STEP 1: PREPROCESS DATA .....</b>	<b>4</b>
<b>STEP 2: CHOOSE THE ARCHITECTURE .....</b>	<b>5</b>
ONE HIDDEN LAYER .....	5
<i>Compare regularization effect .....</i>	<i>6</i>
<i>Learning rate.....</i>	<i>6</i>
TWO HIDDEN LAYERS .....	8
<i>50 and 25 nodes.....</i>	<i>9</i>
<i>2000 and 1000 nodes.....</i>	<i>10</i>

**Dataset:** <http://mmlab.ie.cuhk.edu.hk/projects/CelebA.html>

The dataset contains more than 200K images with its respective bounding box face annotation.

#### **Training Dataset:**

**Face:** 10000 images from the dataset are used for training the model. After cropping out the bounding box such that only faces are kept, the images are resized to 60 x 60 x 3-pixel size.

**Non-face:** 10000 images for training are formed by cropping out 60 x 60 x 3 regions from the background of the original training face images.

#### **Test Dataset:**

**Face:** 1000 images from the dataset are used for testing the model. After cropping out the bounding box such that only faces are kept, the images are resized to 60 x 60 x 3-pixel size.

**Non-face:** 1000 images for testing are formed by cropping out 60 x 60 x 3 regions from the background of the original testing face images.

#### **Validation Dataset:**

**Face:** 1000 images from the dataset are used for validating the model. After cropping out the bounding box such that only faces are kept, the images are resized to 60 x 60 x 3-pixel size.

**Non-face:** 1000 images for validation are formed by cropping out 60 x 60 x 3 regions from the background of the original validation face images.

```
26         img=cv2.imread(l[0])
27         x_1=int(l[1])
28         y_1=int(l[2])
29         w=int(l[3])
30         h=int(l[4])
31         #print (l[0])
32         crop_img=img[y_1:y_1+h,x_1:x_1+w]
33         resized_image = cv2.resize(crop_img, (60, 60),interpolation = cv2.INTER_AREA)
34 #         #Set path to the folder where you want to write the cropped face images
35         os.chdir(path/"face")
36         cv2.imwrite(l[0], resized_image)
37         os.chdir(path/"nonface")
38         newimg=img[0:60,0:60,:]
39         cv2.imwrite(l[0],newimg)
40
```

Fig 1. Preparing Data

## Step 1: Preprocess Data

Perform preprocessing on the data to get zero-centered data and normalized data

```
113|  
114| """  
115| Preprocess data to get zero centered and normalized data  
116|  
117| """  
118| scaler = StandardScaler()  
119| scaler.fit(train_data)  
120| StandardScaler(copy=True, with_mean=True, with_std=True)  
121| train_data = scaler.transform(train_data)  
122| test_data = scaler.transform(test_data)  
123| valid_data = scaler.transform(valid_data)  
124|
```

Fig 2. Preprocess data

Training and Testing Accuracy was observed for both preprocessed and non-preprocessed data. There was an increase in accuracy for preprocessed data. The results were compared for a neural network of single hidden layer of 1024 nodes and training batch size of 128. Learning rate of 0.001 and regularization factor (lambda) of 0.01.

Minibatch loss at step 3000: 47315.41796875

Minibatch accuracy: 89.8

Validation accuracy: 83.3

Test accuracy: 83.5

(a) No preprocessing

Minibatch loss at step 3000: 6768.2822265625

Minibatch accuracy: 96.9

Validation accuracy: 85.0

Test accuracy: 85.5

(b) With preprocessing

Fig 3. Compare accuracy

## Step 2: Choose the Architecture

Initially with a simple neural network of 1 Hidden layer only with 1000 nodes and training batch size of 100

```
99     tf_train_dataset = tf.placeholder('float', shape=(batch_size, size))
100     tf_train_labels = tf.placeholder('float', shape=(batch_size, num_labels))
101     tf_valid_dataset = tf.constant(valid_data)
102     tf_test_dataset = tf.constant(test_data)
103
104     # Variables.
105     weights_1 = tf.Variable(tf.truncated_normal([size, num_nodes]))
106     biases_1 = tf.Variable(tf.zeros([num_nodes]))
107     weights_2 = tf.Variable(tf.truncated_normal([num_nodes, num_labels]))
108     biases_2 = tf.Variable(tf.zeros([num_labels]))
109
```

Fig 4. Define the network weights

## Babysitting process

```
109
110 # Training computation.
111 logits_1 = tf.matmul(tf_train_dataset, weights_1) + biases_1
112 relu_layer= tf.nn.relu(logits_1)
113 logits_2 = tf.matmul(relu_layer, weights_2) + biases_2
114 # Normal loss function
115 loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits_2, labels=tf_train_labels))
116 optimizer = tf.train.AdamOptimizer().minimize(loss)
117
118 # Predictions for the training
119 train_prediction = tf.nn.softmax(logits_2)
120
121 # Predictions for validation
122 logits_1 = tf.matmul(tf_valid_dataset, weights_1) + biases_1
123 relu_layer= tf.nn.relu(logits_1)
124 logits_2 = tf.matmul(relu_layer, weights_2) + biases_2
125
126 valid_prediction = tf.nn.softmax(logits_2)
127
128 # Predictions for test
129 logits_1 = tf.matmul(tf_test_dataset, weights_1) + biases_1
130 relu_layer= tf.nn.relu(logits_1)
131 logits_2 = tf.matmul(relu_layer, weights_2) + biases_2
132
133 test_prediction = tf.nn.softmax(logits_2)
134
```

Throughout the project,  
Activation function used was ReLU and Adam Optimizer and loss calculated using softmax cross entropy

Fig 5. Predictions computations

## Compare regularization effect

Initialized	Initialized
Minibatch loss at step 0: 293.8826904296875	Minibatch loss at step 0: 4176613632.0
Minibatch accuracy: 75.0	Minibatch accuracy: 58.0
Validation accuracy: 51.4	Validation accuracy: 48.7
Minibatch loss at step 3000: 3.5568065643310547	Minibatch loss at step 3000: 2339399.25
Minibatch accuracy: 99.0	Minibatch accuracy: 67.0
Validation accuracy: 97.4	Validation accuracy: 71.5
Test accuracy: 98.3	Test accuracy: 72.2

(a) No regularization

(b) Lambda = 1e3 for regularization factor

Fig 6. Regularization, Learning rate was 0.001

Loss went up as regularization factor was increased. Also there is quite a drop in accuracy.

## Learning Rate

Learning rate was changed and accuracy was observed.

Minibatch loss at step 3000: 8046.59521484375	Minibatch loss at step 3000: 32386.2890625
Minibatch accuracy: 96.0	Minibatch accuracy: 97.0
Validation accuracy: 96.2	Validation accuracy: 68.0
Test accuracy: 96.6	Test accuracy: 68.0

(a) Learning Rate = 0.001

(b) Learning Rate = 0.01

Minibatch loss at step 3000: 2756.4150390625	Minibatch loss at step 3000: 25055.68359375
Minibatch accuracy: 99.0	Minibatch accuracy: 94.0
Validation accuracy: 97.5	Validation accuracy: 96.2
Test accuracy: 97.6	Test accuracy: 97.1

(c) Learning Rate = 0.001

(d) Learning Rate = 0.01

Fig 7. Regularization factor (lambda) = 0.01 for (a) and (b)  
Regularization factor (lambda) = 0.001 for (c) and (d)

Learning rate starting at 0.5 with decay rate of 0.96 these observations were obtained

Learning rate: 0.500000  
Minibatch loss at step 0: 43007.8984375  
Minibatch accuracy: 17.0  
Validation accuracy: 52.1  
Learning rate: 0.480000  
Minibatch loss at step 500: 13787915.0  
Minibatch accuracy: 46.0  
Validation accuracy: 64.7  
Learning rate: 0.460800  
Minibatch loss at step 1000: 30334266.0  
Minibatch accuracy: 52.0  
Validation accuracy: 84.8  
Learning rate: 0.442368  
Minibatch loss at step 1500: 2638541.0  
Minibatch accuracy: 95.0  
Validation accuracy: 78.4  
Learning rate: 0.424673  
Minibatch loss at step 2000: 2134269.25  
Minibatch accuracy: 94.0  
Validation accuracy: 79.3  
Learning rate: 0.407686  
Minibatch loss at step 2500: 844442.5  
Minibatch accuracy: 98.0  
Validation accuracy: 77.2  
Learning rate: 0.391379  
Minibatch loss at step 3000: 332023.3125  
Minibatch accuracy: 99.0  
Validation accuracy: 61.8  
Test accuracy: 61.2

With a high learning rate the accuracy is low and loss is high  
It is observed Learning Rate of 0.001 gives best accuracy among these values.

For this neural network architecture of 1 hidden layer highest test accuracy obtained was 98.3% for no regularization and learning rate =0.001.




Fig 8. One Hidden Layer with 1000 nodes  
Test accuracy = 98.3%

## For two Hidden Layers Neural Network

```
159
160 # Variables.
161 weights_1 = tf.Variable(tf.truncated_normal([size,hidden_nodes_1 ]))
162 biases_1 = tf.Variable(tf.zeros([hidden_nodes_1]))
163 weights_2 = tf.Variable(tf.truncated_normal([hidden_nodes_1, hidden_nodes_2]))
164 biases_2 = tf.Variable(tf.zeros([hidden_nodes_2]))
165 # Output layer
166 weights_3 = tf.Variable(tf.truncated_normal([hidden_nodes_2, num_labels]))
167 biases_3 = tf.Variable(tf.zeros([num_labels]))
168
169
170 # Hidden RELU Layer 1
171 logits_1 = tf.matmul(tf_train_dataset, weights_1) + biases_1
172 hidden_layer_1 = tf.nn.relu(logits_1)
173
174
175 # Hidden RELU Layer 2
176 logits_2 = tf.matmul(hidden_layer_1, weights_2) + biases_2
177 hidden_layer_2 = tf.nn.relu(logits_2)
178
179
180 # Output Layer
181 logits_3 = tf.matmul(hidden_layer_2, weights_3) + biases_3
182
183 # Normal loss function
184 loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits_3, labels=tf_train_labels))
185 # Loss function with L2 Regularization
186 regularizers = tf.nn.l2_loss(weights_1) + tf.nn.l2_loss(weights_2) + \
187               tf.nn.l2_loss(weights_3)
188 loss = tf.reduce_mean(loss + beta * regularizers)
189
190 optimizer = tf.train.AdamOptimizer().minimize(loss)
```

Fig 9. Define network weight and architecture for 2 Hidden layers

```
191
192 # Predictions for the training
193 train_prediction = tf.nn.softmax(logits_3)
194
195 # Predictions for validation
196 valid_logits_1 = tf.matmul(tf_valid_dataset, weights_1) + biases_1
197 valid_relu_1 = tf.nn.relu(valid_logits_1)
198
199 valid_logits_2 = tf.matmul(valid_relu_1, weights_2) + biases_2
200 valid_relu_2 = tf.nn.relu(valid_logits_2)
201
202 valid_logits_3 = tf.matmul(valid_relu_2, weights_3) + biases_3
203
204
205 valid_prediction = tf.nn.softmax(valid_logits_3)
206
207 # Predictions for test
208 test_logits_1 = tf.matmul(tf_test_dataset, weights_1) + biases_1
209 test_relu_1 = tf.nn.relu(test_logits_1)
210
211 test_logits_2 = tf.matmul(test_relu_1, weights_2) + biases_2
212 test_relu_2 = tf.nn.relu(test_logits_2)
213
214 test_logits_3 = tf.matmul(test_relu_2, weights_3) + biases_3
215
216
217 test_prediction = tf.nn.softmax(test_logits_3)
218
```

Fig 10. Predictions computations



A second simple neural network of 2 Hidden layers with 50 and 25 nodes and training batch size of 100 was trained and results were observed for Learning Rate = 0.001 and Regularization factor ( $\lambda$ ) = 0.01 Testing accuracy of 95 % was obtained.

Initialized

Minibatch loss at step 0: 2115.986328125

Minibatch accuracy: 97.0

Validation accuracy: 50.0

Minibatch loss at step 200: 2317.554931640625

Minibatch accuracy: 64.0

Validation accuracy: 80.7

Minibatch loss at step 400: 2222.820556640625

Minibatch accuracy: 62.0

Validation accuracy: 82.2

Minibatch loss at step 600: 2453.8408203125

Minibatch accuracy: 71.0

Validation accuracy: 81.0

Minibatch loss at step 800: 1992.5531005859375

Minibatch accuracy: 79.0

Validation accuracy: 85.6

Minibatch loss at step 1000: 1840.5462646484375

Minibatch accuracy: 74.0

Validation accuracy: 88.3

Minibatch loss at step 1200: 1644.425537109375

Minibatch accuracy: 85.0

Validation accuracy: 90.8

Minibatch loss at step 1400: 1592.3096923828125

Minibatch accuracy: 84.0

Validation accuracy: 91.8

Minibatch loss at step 1600: 1519.7130126953125

Minibatch accuracy: 90.0

Validation accuracy: 92.0

Minibatch loss at step 1800: 1452.201904296875

Minibatch accuracy: 84.0

Validation accuracy: 92.8

Minibatch loss at step 2000: 1368.598388671875

Minibatch accuracy: 90.0

Validation accuracy: 92.8

Minibatch loss at step 2200: 1288.3709716796875

Minibatch accuracy: 94.0

Validation accuracy: 92.1

Minibatch loss at step 2400: 1226.506591796875

Minibatch accuracy: 85.0

Validation accuracy: 93.7

Minibatch loss at step 2600: 1146.2833251953125

Minibatch accuracy: 93.0

Validation accuracy: 93.8

Minibatch loss at step 2800: 1076.744384765625

Minibatch accuracy: 90.0

Validation accuracy: 94.1

Minibatch loss at step 3000: 1003.8374633789062

Minibatch accuracy: 84.0

Validation accuracy: 94.0

Test accuracy: 95.0




Fig 11. Two Hidden Layers with 50 and 25 nodes  
Test accuracy = 95 %

A third simple neural network of 2 Hidden layers with 2000 and 1000 nodes and training batch size of 100 was trained and results were observed for Learning Rate = 0.001 and Regularization factor ( $\lambda$ ) = 0.01 Testing accuracy of 98.1 % was obtained.

Initialized

Minibatch loss at step 0: 148128.9375  
Minibatch accuracy: 4.0  
Validation accuracy: 51.4  
Minibatch loss at step 200: 421059.84375  
Minibatch accuracy: 17.0  
Validation accuracy: 63.9  
Minibatch loss at step 400: 281575.96875  
Minibatch accuracy: 52.0  
Validation accuracy: 76.7  
Minibatch loss at step 600: 322211.6875  
Minibatch accuracy: 52.0  
Validation accuracy: 75.2  
Minibatch loss at step 800: 231896.703125  
Minibatch accuracy: 63.0  
Validation accuracy: 81.6  
Minibatch loss at step 1000: 221773.609375  
Minibatch accuracy: 74.0  
Validation accuracy: 84.2  
Minibatch loss at step 1200: 155260.5625  
Minibatch accuracy: 71.0  
Validation accuracy: 88.7  
Minibatch loss at step 1400: 122696.7421875  
Minibatch accuracy: 84.0  
Validation accuracy: 91.7  
Minibatch loss at step 1600: 95274.1796875  
Minibatch accuracy: 91.0  
Validation accuracy: 94.2  
Minibatch loss at step 1800: 93604.8359375

Minibatch accuracy: 93.0  
 Validation accuracy: 95.8  
 Minibatch loss at step 2000: 89518.3515625  
 Minibatch accuracy: 96.0  
 Validation accuracy: 96.3  
 Minibatch loss at step 2200: 87143.0859375  
 Minibatch accuracy: 98.0  
 Validation accuracy: 97.0  
 Minibatch loss at step 2400: 85454.53125  
 Minibatch accuracy: 98.0  
 Validation accuracy: 97.3  
 Minibatch loss at step 2600: 86446.703125  
 Minibatch accuracy: 97.0  
 Validation accuracy: 97.3  
 Minibatch loss at step 2800: 84319.6015625  
 Minibatch accuracy: 97.0  
 Validation accuracy: 97.4  
 Minibatch loss at step 3000: 82522.2734375  
 Minibatch accuracy: 99.0  
 Validation accuracy: 97.5  
 Test accuracy: 98.1




Fig 12. Two Hidden Layers with 2000 and 1000 nodes  
 Test accuracy = 98.1%